

COSC 411: Artificial Intelligence

Uninformed Search

Dr. Shuangquan (Peter) Wang
(spwang@salisbury.edu)

Department of Computer Science
Salisbury University




About this note

- Most slides of this note are from:
 - "CS 188 Introduction to Artificial Intelligence" teaching material, UC Berkeley, Summer 2020 (instructor: Nikita Kitaev)
 - <https://inst.eecs.berkeley.edu/~cs188/su20/>
 - "CS 188 Introduction to Artificial Intelligence" teaching material, UC Berkeley, Spring 2019 (instructor: Sergey Levine and Stuart Russell)
 - <https://inst.eecs.berkeley.edu/~cs188/sp19/>


Dissemination or sale of any part of this note is NOT permitted!

Contents

- Agents that Plan Ahead
- Search Problems
- Uninformed Search Methods
 - Depth-First Search
 - Breadth-First Search
 - Uniform-Cost Search



Agents and environments



- An agent **perceives** its environment through **sensors** and **acts** upon it through **actuators**


Rationality

- A **rational agent** chooses actions maximize the **expected** utility
 - agents that have a goal, and a cost
 - E.g., reach goal with lowest cost
 - agents that have numerical utilities, rewards, etc.
 - E.g., take actions that maximize total reward over time (e.g., largest profit in \$)

Agent design


- The environment type largely determines the agent design
 - *Fully/partially observable* => agent requires *memory* (internal state)
 - *Discrete/continuous* => agent may not be able to enumerate *all states*
 - *Stochastic/deterministic* => agent may have to prepare for *contingencies*
 - *Single-agent/multi-agent* => agent may need to behave *randomly*

Agents that Plan



Reflex Agents


- Reflex agents:
 - Choose action based on current percept (and maybe memory)
 - May have memory or a model of the world's current state
 - Do not consider the future consequences of their actions
 - Consider how the world **IS**
- Can a reflex agent be rational?



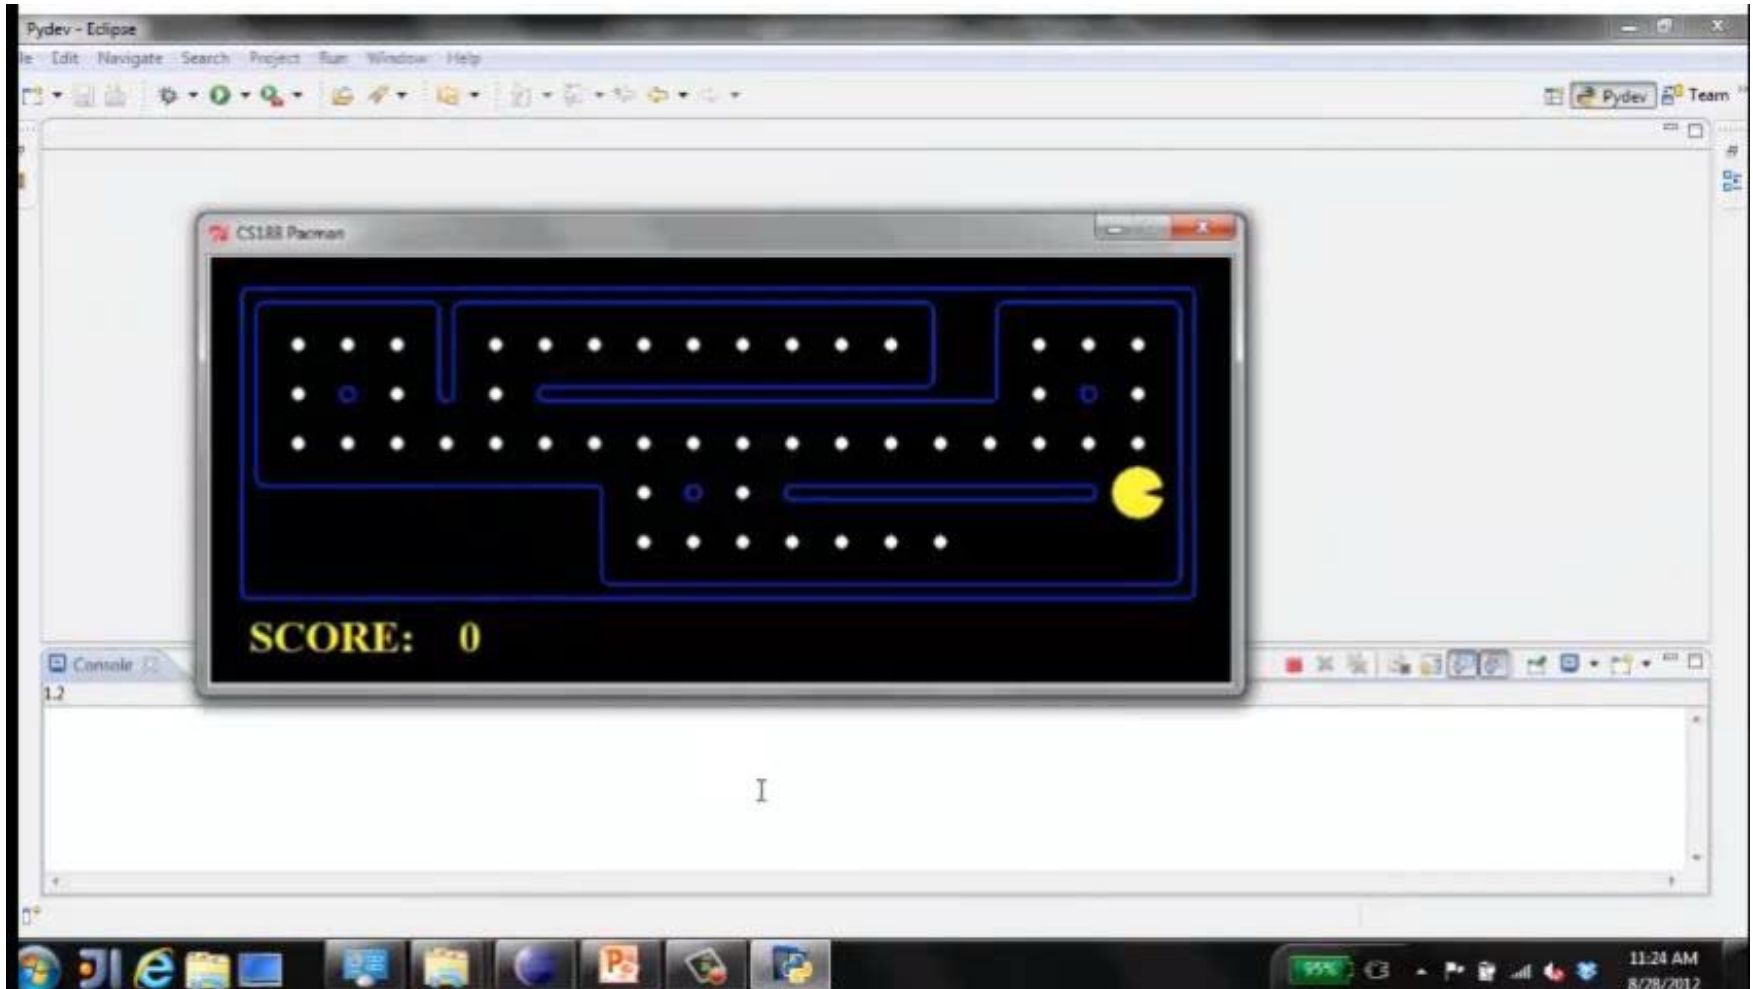
Planning Agents

- Planning agents:

- Ask “what if”
- Decisions based on (hypothesized) consequences of actions
- Must have a model of how the world evolves in response to actions
- Must formulate a goal (test)
- Consider how the world **WOULD BE**




Search Problems (Pac-Man)




Search Problems

- A **search problem** consists of:

- A state space



- A successor function
(with actions, costs)



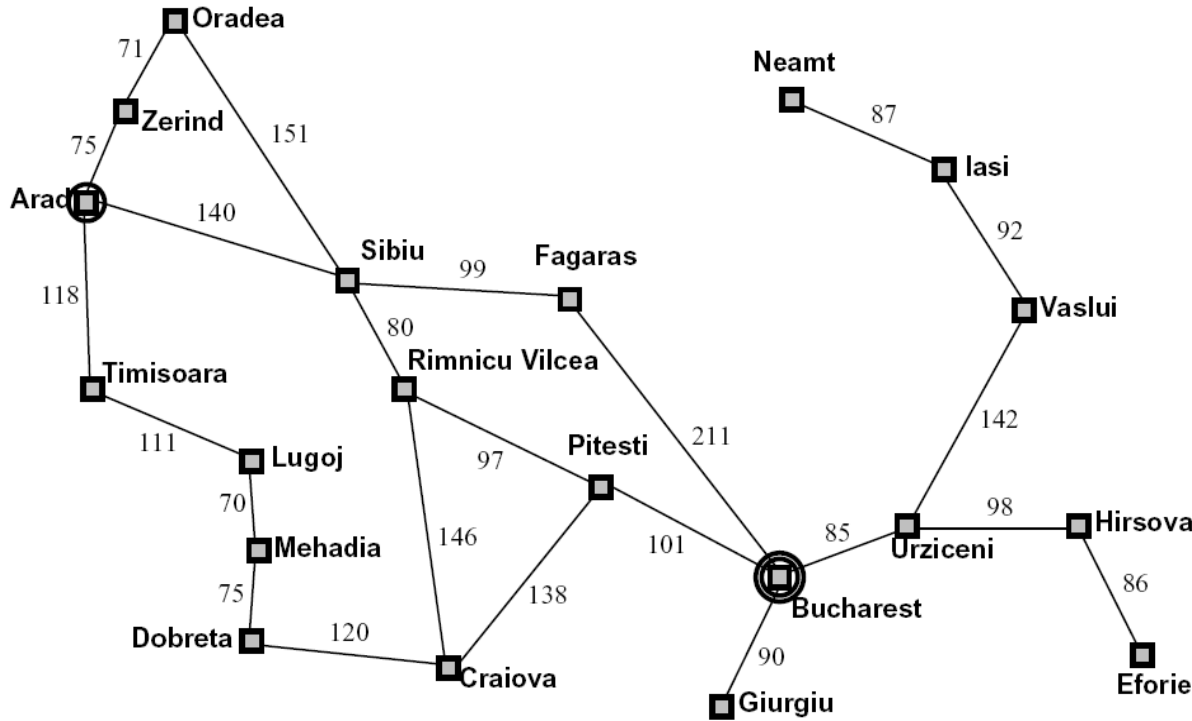
- A start state and a goal test

- A **solution** is a sequence of actions (a plan) which transforms the start state to a goal state

Search Problems Are Models




Example: Traveling in Romania



- State space:
 - Cities
- Successor function:
 - Roads: Go to adjacent city with cost = distance
- Start state:
 - Arad
- Goal test:
 - Is state == Bucharest?
- Solution?

What's in a State Space?

The **world state** includes every last detail of the environment



A **search state** keeps only the details needed for planning (abstraction)

■ Problem: Pathing


- States: (x,y) location
- Actions: N/S/E/W
- Successor: update location only
- Goal test: is $(x,y)=\text{END}$

■ Problem: Eat-All-Dots


- States: $\{(x,y), \text{dot booleans}\}$
- Actions: N/S/E/W
- Successor: update location and possibly a dot boolean
- Goal test: dots all false

State Space Sizes?

- World state:
 - Agent positions: 120
 - Food count: 30
 - Ghost positions: 12
 - Agent facing: N/S/E/W
- How many
 - World states?
 $120 \times (2^{30}) \times (12^2) \times 4$
 - States for pathing?
120
 - States for eat-all-dots?
 $120 \times (2^{30})$




State Space Graphs and Search Trees




State Space Graphs

- State space graph: A mathematical representation of a search problem
 - Nodes are (abstracted) world configurations
 - Arcs represent successors (action results)
 - The goal test is a set of goal nodes (maybe only one)
- In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea



Search Trees




- A search tree:


- A “what if” tree of plans and their outcomes
- The start state is the root node
- Children correspond to successors
- Nodes show states, but correspond to PLANS that achieve those states
- For most problems, we can never actually build the whole tree

State Space Graphs vs. Search Trees

State Space Graph




Search Tree




*Each NODE in the search tree is an entire PATH in the state space graph.
We construct both on demand – and we construct as little as possible.*

State Space Graphs vs. Search Trees

Consider this 4-state graph:




How big is its search tree (from S)?

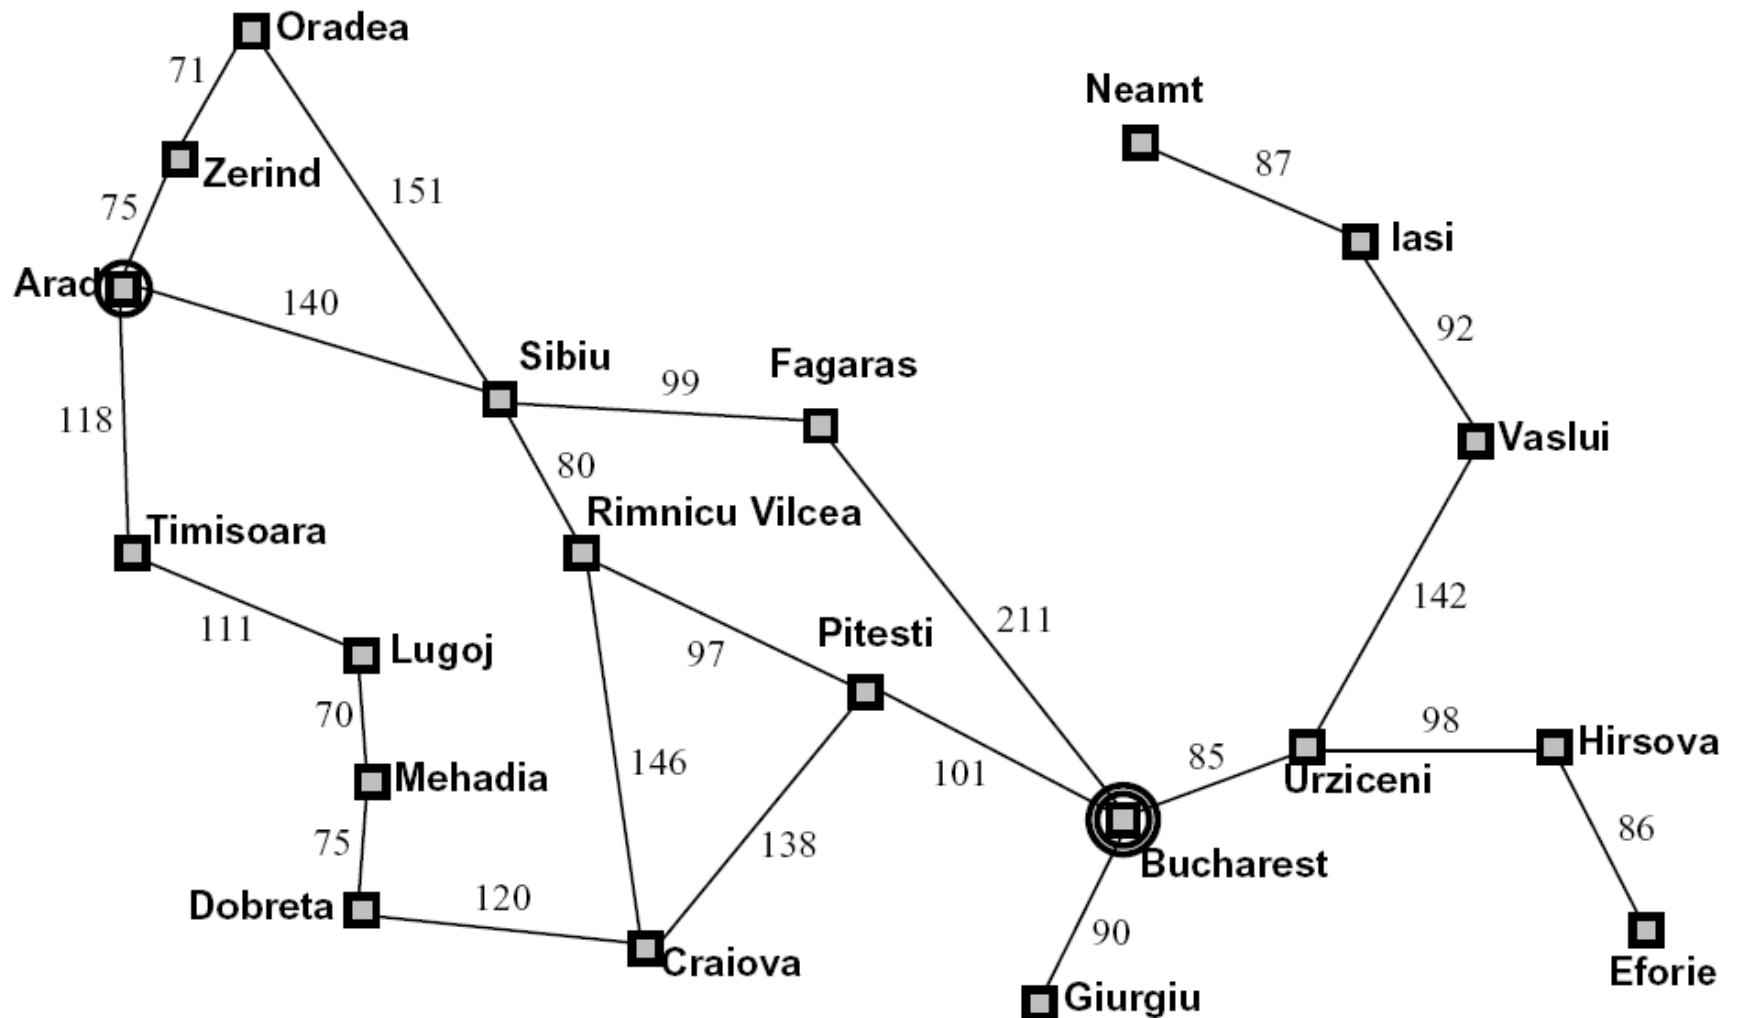


Important: Lots of repeated structure in the search tree!

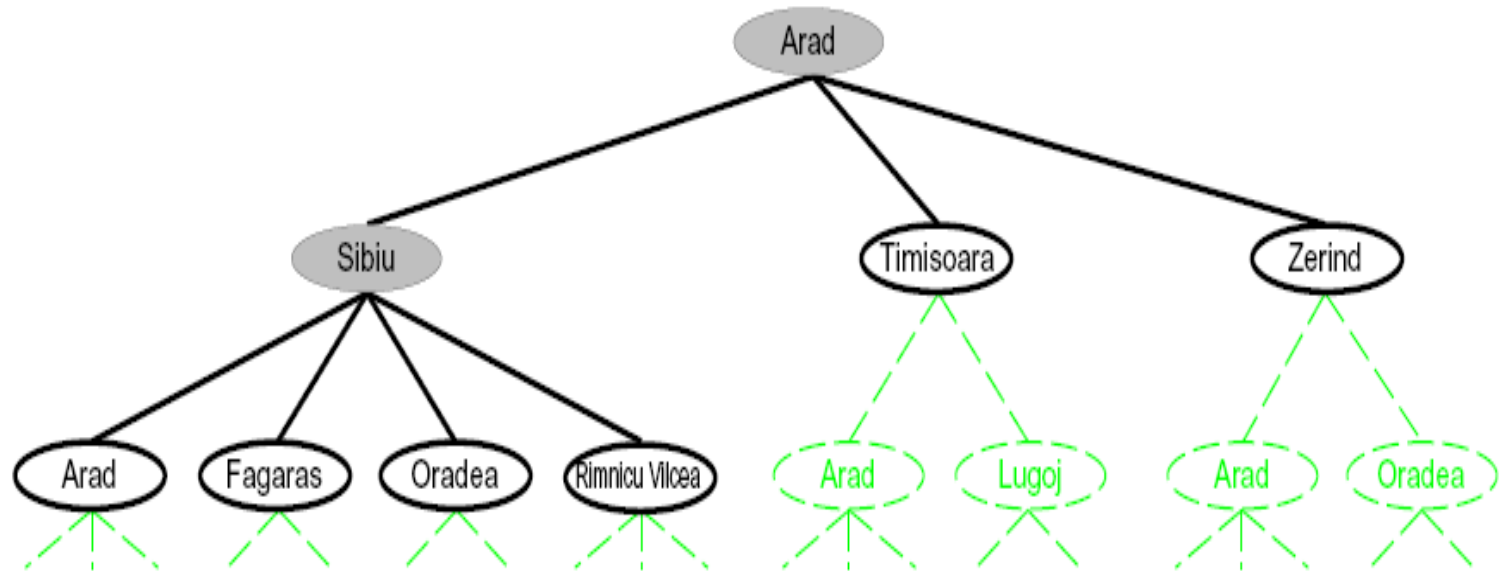
Tree Search



Search Example: Romania



Searching with a Search Tree



- Search:
 - Expand out potential plans (tree nodes)
 - Maintain a **fringe** of partial plans under consideration
 - Try to expand as few tree nodes as possible


General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```


- Important ideas:
 - Fringe
 - Expansion
 - Exploration strategy
- Main question: which fringe nodes to explore?

Tree Search


Breadth
First
Search



Depth
First
Search





Depth-First Search



Depth-First Search

Strategy: expand a deepest node first

*Implementation:
Fringe is a **LIFO stack***



Search Algorithm Properties


- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?

- Cartoon of search tree:

- b is the branching factor
- m is the maximum depth
- solutions at various depths


- Number of nodes in entire tree?

- $1 + b + b^2 + \dots + b^m = O(b^m)$




Depth-First Search (DFS) Properties

- What nodes DFS expand?
 - Some left prefix of the tree.
 - Could process the whole tree!
 - If m is finite, takes time $O(b^m)$
- Is it complete?
 - m could be infinite, so only if we prevent cycles (more later)
- Is it optimal?
 - No, it finds the “leftmost” solution, regardless of depth or cost





Breadth-First Search



Breadth-First Search


Strategy: expand a shallowest node first

*Implementation: Fringe is a **FIFO queue***




Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
 - Processes all nodes above shallowest solution
 - Let depth of shallowest solution be s
 - Search takes time $O(b^s)$
- Is it complete?
 - s must be finite if a solution exists, so yes!
- Is it optimal?
 - Only if costs are all 1 (more on costs later)




Quiz: DFS vs BFS




- When will BFS outperform DFS?
- When will DFS outperform BFS?

Cost-Sensitive Search



BFS finds the shortest path in terms of number of actions.
It does not find the least-cost path. We will now cover
a similar algorithm which does find the least-cost path.



Uniform Cost Search



Uniform Cost Search


Strategy: expand a
cheapest node first:

Fringe is a priority queue
(priority: cumulative cost)





Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
 - Processes all nodes with cost less than cheapest solution!
 - If that solution costs C^* and arcs cost at least ε , then the “effective depth” is roughly C^*/ε
 - Takes time $O(b^{C^*/\varepsilon})$ (exponential in effective depth)
- Is it complete?
 - Assuming best solution has a finite cost and minimum arc cost is positive, yes!
- Is it optimal?
 - Yes!




Uniform Cost Issues


- Remember: UCS explores increasing cost contours
- The good: UCS is complete and optimal!
- The bad:
 - Explores options in every “direction”
 - No information about goal location




Video of Demo Empty UCS




Video of Demo Maze with Deep/Shallow Water --- DFS, BFS, or UCS? (part 1)



Video of Demo Maze with Deep/Shallow Water - -- DFS, BFS, or UCS? (part 2)




Video of Demo Maze with Deep/Shallow Water --- DFS, BFS, or UCS? (part 3)




Search and Models

- Search operates over models of the world
 - The agent doesn't actually try all the plans out in the real world!
 - Planning is all “in simulation”
 - Your search is only as good as your models...



Search Gone Wrong?



Thanks